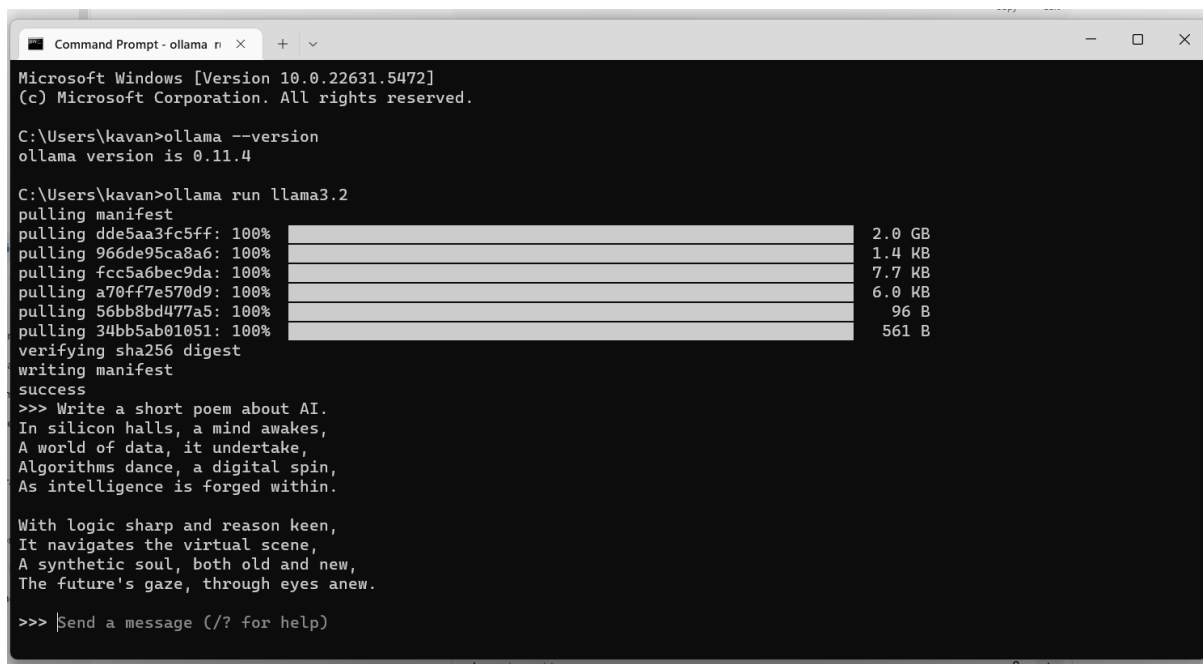


Question 4 (ollama)

Result:



```
Microsoft Windows [Version 10.0.22631.5472]
(c) Microsoft Corporation. All rights reserved.

C:\Users\kavan>ollama --version
ollama version is 0.11.4

C:\Users\kavan>ollama run llama3.2
pulling manifest
pulling dde5aa3fc5ff: 100% 2.0 GB
pulling 966de95ca8a6: 100% 1.4 KB
pulling fcc5a6bec9da: 100% 7.7 KB
pulling a70ff7e570d9: 100% 6.0 KB
pulling 56bb8bd477a5: 100% 96 B
pulling 34bb5ab01051: 100% 561 B
verifying sha256 digest
writing manifest
success
>>> Write a short poem about AI.
In silicon halls, a mind awakes,
A world of data, it undertake,
Algorithms dance, a digital spin,
As intelligence is forged within.

With logic sharp and reason keen,
It navigates the virtual scene,
A synthetic soul, both old and new,
The future's gaze, through eyes anew.

>>> Send a message (? for help)
```

Performance and installation experience:

The installation process was straightforward. I downloaded the Ollama installer from the official website and set it up on Windows without any major issues. After confirming the installation using the `ollama --version` command, I proceeded to run the model `llama3.2`.

The first run required downloading the model weights, which was around 4GB. This step took a few minutes depending on my internet speed. Once the model was downloaded, it launched successfully in the Command Prompt with the `>>>` prompt, ready to accept input.

I tested it by asking:

Prompt: *"Write a short poem about AI."*

The model responded within a few seconds, generating a short, creative poem. The response time after the first load was quite fast (about 3–5 seconds).

Overall, the installation experience was positive. Ollama made it easy to run a local LLM without complicated setup steps, and the model's response quality was impressive for a fully local environment.

Question 4 (ollama)